

The RAG demo system uses PostgreSQL with pgvector for vector storage.

Embeddings are generated using the Hugging Face inference API.

Cosine similarity ranks the most relevant document chunks.